



Top-Down vs Bottom-Up Design

Common pitfall: start with the details first



Top-Down vs Bottom-Up Design

Better solution: Figure out how all the components fit together first, then build them



What is an API?

Application **P**rogramming **I**nterface

Sits between layers or components in a system

Defines a **contract** for how the components will behave

Anything that an external component **relies on** is part of the API

API Example

Cars have an API:

- Use the steering wheel to steer, the gas to go, the brake to stop

Users don't need to know the details:

- Drive-by-wire vs mechanical linkages
- Different types of brake disks
- How an internal combustion engine works

Why APIs?

Allows **interacting** with a component without knowing the **details** of how it works

Important for scaling software:

- Each piece is complicated, no one can keep track of all details at once

- APIs let you focus on just one piece at a time

Any separation between components has an **implicit** API, but not necessarily a good one (organic APIs are usually bad)

Why APIs?

Side benefit: Promotes better design


Highly-coupled components tend to be **fragile**

APIs push design to a **top-down** approach, which produces more flexible and simpler coding than bottom-up design

Component APIs

Any boundary between components needs an API

- **network** boundaries
- **process** boundaries
- **conceptual** boundaries



System Design Diagram


Used to identify where APIs are needed

Helpful first step before architecting a large system

Useful for parallelizing work across many people/teams

Acts like an informal UML

Not a Flowchart



Systems Design Diagrams

Start with an overall workflow for the system

1. **Initial Layout:** For each step, assign it to a component - draw a **box** and label it
2. **Environmental Features:** Surround any components that live on either the same physical **server** or in the same **process** with an extra box
3. **Interactions:** Each time a component interacts with another, draw an **arrow** to connect them - these will be the APIs
4. **Simplify:** Sometimes two components will be the **request/response** steps for a single user action; identify these and combine them into a single component. Similarly, if any components are the **only component** within a server/process boundary, you can remove the double-box around them.

Note: a component from 'outside' a box should only interact with the outermost layer, which in turn will interact with internal components; this separates external and internal APIs to simplify backwards-compatibility guarantees

Systems Design Diagrams

When have you simplified enough?

- Whenever possible, have **pairs of arrows** between components representing **request/response** pairs

A system with a **single arrow** is unusual and more complicated:

- A single arrow means there **is no response** to indicate that a task is completed, successful, or failed
- Any such system also needs components to **monitor and potentially redo** the task in question

Systems Design Diagrams - Removing Single Arrows

Protocols with no response (SMTP): "resend this email"

- No way to remove the single arrow, just build a way to re-send

Tasks with delayed response:

- Add some sort of id/lookup key response
- Add another component to "check status" using the key

ex: placing an order on a website returns an order number, which allows someone to look up the status of the order over several days

Systems Diagrams Example

Example: Adding a 'Search' feature to an application

Overall workflow:

A user will enter terms in the search bar and hit enter. That request will go to a server that will create a parsed version of the query, then compare that against indexed data. The results will be compiled into a user-friendly format and then returned to the application.

Systems Diagrams Example

Overall workflow:

A user will enter terms in the search bar and hit enter. That request will go to a server that will create a parsed version of the query, then compare that against indexed data. The results will be compiled into a user-friendly format and then returned to the application.

Step 1: Create a component for each step



Systems Diagrams Example

Step 2: Add server/process boundary boxes

Search Bar


**Create Parsed
Query**

**Compare to
Indexed Data**

**Compile Results
and Return to
User**


Systems Diagrams Example

Step 3: Draw component interactions



Systems Diagrams Example


Step 4: Combine request/response components & remove extra boxes



Systems Diagrams Example

Now we can read the required APIs off the diagram:

- one network API from the search bar to the network-facing server (needs **wire** compatibility)
- one internal/in-process API from the networking layer to the query-parsing/translation component
- one internal/in-process API from the parsing component to the indexed data searching component.



Your turn!

Create a systems diagram for the following workflow:

A user will enter data in a website form and submit it. That information will get sent to the server, which will transform it into a SQL query and send that to a database. The database results will get transformed back into javascript-friendly results, and returned to the website to be displayed.

Run through the four steps to create the systems diagram:

1. **Initial Layout:** For each step, assign it to a component - draw a **box** and label it
2. **Environmental Features:** Surround any components that live on either the same physical **server** or in the same **process** with an extra box
3. **Interactions:** Each time a component interacts with another, draw an **arrow** to connect them - these will be the APIs
4. **Simplify:** Sometimes two components will be the **request/response** steps for a single user action; identify these and combine them into a single component. Similarly, if any components are the **only component** within a server/process boundary, you can remove the double-box around them.

Web Form Systems Diagram

A user will enter data in a website form and submit it. That information will get sent to the server, which will transform it into a SQL query and send that to a database. The database results will get transformed back into javascript-friendly results, and returned to the website to be displayed.



Web Form Systems Diagram (Alt)

What's different here?


Why would you choose one vs the other?



Designing the API

The API is the layer **between** components, and acts as a **contract** between them

Each component has a copy of the contract, but "the API" isn't a single piece of code



Designing a Good API

APIs should be designed from the perspective of the **user**, not the implementer

- Avoids a "leaky abstraction", where details of how a component works end up in the API
- Captures behavior, not data (although the line can be blurry)
- Needs enough power to be useful, but also an easy learning curve
- Make the common things easy, and the weird things possible

Remember: you can always **add** things to an API, but it's difficult to **remove** or **change** things

Initially, aim for a Minimum Viable Product (MVP) approach - what's the smallest set of things that will let a user accomplish what they need to

Backwards Compatibility

For widely used, public, open source projects (Ex: protocol buffers, the JRE):

- Many clients
- No complete list of all clients, much less access to their code

⇒ Breaking changes are a Big Deal, and very rare

- Ever wondered why generics are wonky in Java? To maintain 1.4/1.5 compatibility

Backwards Compatibility

Some APIs have a narrower scope: within a company, team, or project

Backwards compatibility is still important!

Consider a simple web server: a client process and a server process


When upgrading, one option is to turn everything off, upgrade everything, and then turn it back on

- Ex: "Caution: AmateurHour.com will be going down for maintenance from 3am to 6am on Friday"

Online Updates


Better option: At least 2 copies of every process with network failover.

Ex: a user is connected to Client 1, which starts off connected to Server 1; both clients/servers are running identical software




Online Updates

To update to v1.1: first stop the Server 1 process. Client 1 will fail over to Server 2; this will at worst cause a browser-refresh style hiccup for a user, but typically it's not noticeable




Online Updates

Now, update the binaries on the server, start Server 1 back up, and do the same for Server 2




Online Updates

Once the servers are updated, the clients can get sequentially updated as well. All of this relies on a client running v1.0 still being able to talk to a server with v1.1, which means the API must be **backwards compatible** for at least one version



Online Updates

In practice, this whole process is often automated, with Server 1 having to pass some automated health checks before Server 2 gets upgraded, as part of a **continuous deployment** pipeline



Types of Backwards Compatibility

APIs are contract between different components, which may change and upgrade at different rates, with different people

It's important to maintain **backwards compatibility** with an API

What type of backward compatibility is actually **part** of the API:

- Source compatibility: code written with a previous version still compiles
- Wire compatibility: an application built against a previous version can still talk to a new server
- Semantic compatibility: an application's behavior doesn't change

Examples of Breaking Changes

Source compatibility: code written with a previous version still compiles

Previous API:

```
public Set<Object> createCollection();
```

Breaking Change:

```
public Collection<Object> createCollection();
```

Non-breaking Change: returning a HashSet instead of a TreeSet

⇒ Make your return types as **general** as possible

Examples of Breaking Changes

Wire compatibility: an application built against a previous version can still talk to a new server

Previous API:

Using default Java serialization:

```
public Widget createWidget();
```

Breaking Change:

Create a new class that implements `Widget` and return it; Java serialization will fail on the application side

Non-breaking Change:

Use an explicitly backwards/forwards compatible wire format such as JSON or protocol buffers

⇒ Make **serialization** and **persistence** explicit parts of the API

Design For Change - Versioning

Any time something is **serialized** or **persisted** it should have a version number

- Default for java Serializable is serialVersionUID

Two major properties:

- The version number should an incrementing sequence
- Version should be automatically set given the compiled version of the code:
ex: `protocolBufferBuilder.setVersion(VERSION);`

Examples of Breaking Changes

Semantic compatibility: an application's behavior doesn't change

Previous API:

```
public float parseString(String number)
throws NumberFormatException;
```

Breaking Change:

```
/* Returns Float.NaN if the number
doesn't parse */
public float parseString(String
number);
```

⇒ Build in **error handling** from the beginning

Non-breaking Change:

```
@deprecated /** Prefer parseStringNoThrow */
public default float parseString(String
number) throws NumberFormatException {
    float result = parseStringNoThrow(number);
    if (result == Float.NaN) {
        throw new NumberFormatException();
    }
}
```

```
/* Returns Float.NaN if the number doesn't
parse */
public float parseStringNoThrow(String
number);
```

Number of Versions of Compatibility

Large open-source projects will typically support APIs for many years

- Breaking changes are indicated by incrementing the **major** version number

Internal-only or smaller projects may have shorter support windows

- Most common is one **minor** release

Version numbers: major.minor.hotfix/patch (ex: Java 1.6.u45)

Generally, **deprecate** functionality for one support window, then remove it

API Design Feels Backwards

Up until now: I want to build a thing → I go build the thing

New and improved: I want to build a thing → I go do something else entirely, and then eventually come back to building the thing

Seems **counterintuitive**, but trust the process! Somewhere around the middle of the semester, suddenly all that "something else entirely" will really start to pay off

API Design Feels Backwards

Initial API Design steps don't seem to be **getting anywhere**

Think of it like an image loading over time - the first few steps are really blurry, but they're still important



API Example: Building a Server For a Website

After building the System Design Diagram, focus on each API between components

Ex:

A user (via their browser) will connect to a web server for a social media site. Initially, we just want the user to be able to view their profile and make changes to it.

API Example: Building a Server For a Website

Step 1: Identify the **client**

In this case, the website frontend

Note that the client isn't always a person, or the **end user** - the client of an API is the person or (usually) program that will be interacting with it directly



Your Turn! (Part 1 of Many)

After seeing how each piece works with our social media site example, you'll implement the same part of the process with a different system

Goal: You've been asked to build a key-value store system, where another program can store and retrieve arbitrary bytes of data.

Step 1: Who is the **client** for this system?

Part 1 Solution

Goal: You've been asked to build a key-value store system, where clients can store and retrieve arbitrary bytes of data.

The **client** will likely be another piece of code:

- web server
- another backend server
- probably not an actual person

⇒ The user can be trusted (mostly)

⇒ The user (for a given system) is likely **unique** - no need to manage different profiles

API Example: Building a Server For a Website

Step 2: List out what the user needs to be able to do.

In our case, that might be:

- Log in to the site
- Load a profile
- Make changes to a profile



Your Turn! (Part 2)

Goal: You've been asked to build a key-value store system, where clients can store and retrieve arbitrary bytes of data.

The **client** will be another piece of code (some sort of server)

Step 2: List out what operations the client will need to perform

Step 2 Solution

Goal: You've been asked to build a key-value store system, where clients can store and retrieve arbitrary bytes of data.

The functionality it needs to support:

1. Store bytes of data
2. Later, look up the corresponding data

Writing the Actual API

Step 3: Build a prototype

To go from a high-level approach to the actual code, build a **prototype use case**.

This isn't an actual website, or anything that will ever be used; it's to iron out the details to make sure the API is both well-specified and sufficiently powerful

Where the actual website will have forms for input, responses to user actions, etc, the prototype just focuses on the **interactions** with the API

Start with the high-level approach as comments, and fill it out as you go

Writing the Actual API (Next Checkpoint)

Prototypes are frequently **JUnit Tests**, but we won't cover testing until next class.

For now, the prototype for SomeApi will have the format:

PrototypeSomeApi.java:

```
public class PrototypeSomeApi {  
    public void prototype(SomeApi server) {  
        // prototype goes here  
    }  
}
```

This code isn't accessible from any code entry points (no main method or `@Test` annotation), so it will never **run**, that's ok

Prototype code just needs to **compile**

Prototype the Use Case

First, just add comments

```
public void prototype(WebServer server) {  
    // log in the user  
    // load their profile  
    // make a change to the profile  
    // reload the updated version of the profile  
}
```

This step will often flag functionality that was missed during the initial brainstorming, such as a 'log out' option



Prototype the Use Case

First, just add comments

```
public void prototype(WebServer server) {  
    // log in the user  
    // load their profile  
    // make a change to the profile  
    // reload the updated version of the profile  
    // logout  
}
```

All fixed!



Making Code Compile

The previous slide created a new **type** (`WebServer`)

As you're writing the code, you want it to compile - when you have a referenced type that you haven't implemented yet, create it as an **empty interface**:

`WebServer.java`:

```
public interface WebServer {  
    // nothing here yet!  
}
```

⇒ Your IDE can help you with this! Try hovering over the red squiggle, and use the option to create a new type

Your Turn! (Part 3)

Goal: You've been asked to build a key-value store system, where clients can store and retrieve arbitrary bytes of data.

The **client** is another piece of code/other server

The functionality it needs to support:

1. Store bytes of data
2. Later, look up the corresponding data

Create a **PrototypeDataStore.java** with the comments + methods version of the API

Hint: This should compile, so you may need to create an empty interface

Part 3 Solution

```
public void prototype(DataStore dataStore) {  
    // store some data  
    // retrieve the data  
}
```

DataStore.java:

```
public interface DataStore {  
    // nothing here yet  
}
```

Prototype the Use Case

Step 4: Fill out the template

```
public void prototype(WebServer server) {  
    // log in the user: what do we need to build a LoginRequest?  
    LoginResponse response = server.login(new LoginRequest());  
  
    // load their profile  
    // make a change to the profile  
    // reload the updated version of the profile  
    // log out  
}
```

Often, filling in the first few lines doesn't seem like it's actually clarifying anything, but you'll notice we now have two concrete interfaces and a method. Remember, this first pass is still fairly high-level and vague; leave further needed details as comments



API Design Best Practice: Always Have a Return Value

```
LoginResponse login(LoginRequest loginRequest);
```

Why not **void**?

Return values allow you to:

- Handle error cases more gracefully
- Propagate exceptions across network/process boundaries
- Add functionality later without breaking backwards compatibility

Prototype the Use Case

```
public void testWebServer(WebServer server) {  
    // log in the user: what do we need to build a LoginRequest?  
    LoginResponse loginResponse = server.login( new LoginRequest());  
  
    // load their profile  
    if (loginResponse.success()) {  
        ProfileLoadResponse profileLoadResponse =  
server.loadProfile(loginResponse.getUserIdentifier());  
    }  
  
    // make a change to the profile  
    // reload the updated version of the profile  
    // log out  
}
```

As we fill out the next few lines, we'll start to implicitly fill out some of the interfaces; we need to know if the login worked, and in order to load a profile we need a user id, which we would have to get from the login info (good thing we had a response!)

API Design Best Practice: Wrapper Interfaces

```
ProfileLoadResponse loadProfile(UserIdentifier user);
```

Why not 'String username' or 'long userId'?

Exactly!

How the system wants to implement a unique id per user is an **implementation detail**

Leaving it unspecified allows for a flexible implementation, or multiple/changing implementations

Prototype the Use Case

```
public void testWebServer(WebServer server) {  
    // log in the user: what do we need to build a LoginRequest?  
    LoginResponse loginResponse = server.login( new LoginRequest());  
  
    // load their profile  
    if (loginResponse.success()) {  
        ProfileLoadResponse profileLoadResponse =  
server.loadProfile(loginResponse.getUserIdentifier());  
  
        // make a change to the profile: what do we need to create one of these  
        ProfileChangeRequest changeRequest = new ProfileChangeRequest();  
        ProfileChangeResponse profileChangeResponse =  
server.updateProfile(loginResponse.getUserIdentifier(), changeRequest);  
  
        // reload the updated version of the profile  
        profileLoadResponse = server.loadProfile(loginResponse.getUserIdentifier());  
  
        // log out  
        server.logout(loginResponse.getUserIdentifier());  
    }  
}
```

API Design Best Practice: Many Object Interfaces

```
ProfileChangeResponse changeProfile(UserIdentifier user,  
ProfileChangeRequest request);
```

Shouldn't we be reusing code and combining some of these response values?

Nope!

- Copy-pasting **logic** is bad, that's something you encapsulate and re-use
- Reusing **methods** is also a good idea
- On the other hand, allow wrapper **interfaces** to proliferate wildly

Leaving these separate allows for future flexibility

- reuse is great for situations where you control all uses of the code, but APIs are the opposite

tldr: **Types** = don't reuse, **Methods** = do reuse

Writing the actual API: Method Signatures

```
public interface WebServer {  
    LoginResponse login(LoginRequest loginRequest);  
    ProfileLoadResponse loadProfile(UserIdentifier user);  
    ProfileChangeResponse changeProfile(UserIdentifier user,  
ProfileChangeRequest request);  
    LogoutResponse logout(UserIdentifier user);  
}
```

Your Turn!

Recall from earlier:

```
public void prototype(DataStore dataStore) {  
    // store some data  
    // retrieve the data  
}
```

DataStore.java:

```
public interface DataStore {  
    // nothing here yet  
}
```

Now, fill out each commented line, generating interfaces/methods in DataStore as necessary to make everything compile

Prototyped Data Storage Layer

```
public void prototype(DataStore dataStore) {  
    // store some data  
    DataStorageResponse storeResponse =  
dataStore.storeData(new DataStorageRequest());  
  
    // retrieve the data  
    DataLoadResponse loadResponse =  
dataStore.loadData(storeResponse.getDataKey());  
}
```

Prototyped Data Storage Layer

```
public interface DataStore {  
    DataStorageResponse storeData(DataStorageRequest request);  
    DataLoadResponse loadData(DataStorageKey key);  
}
```

Prototyped Data Storage Layer

```
public interface DataStorageRequest {  
    // TBD  
}
```

```
public interface DataStorageResponse {  
    DataStorageKey getDataKey();  
}
```

```
public interface DataStorageKey {  
    // TBD  
}
```

```
public interface DataLoadResponse {  
    // TBD  
}
```

Writing the actual API: Filling in the Rest

```
public interface LoginResponse {  
    boolean success(); // Oh no! Booleans are bad  
    UserIdentifier getUserIdentifier();  
}
```

```
public interface UserIdentifier {  
    // marker interface  
}
```

API Design Best Practice: Booleans are Bad (Usually)

Problem 1: Very inflexible (only ever has two values)

- What happens if we want to add details about why the login failed?

Problem 2: Not self-documenting

- ex: `server.loadProfile(true, false, true, true, false, "");`

The fix: Enums

Enums (Enumerations)

Alternative to listing out a bunch of constants

Use `==`, not `.equals()`

Should only be used for known-at-compile-time, immutable state

Updated API

```
public interface LoginResponse {  
    LoginResponseCode getResponseCode();  
    UserIdentifier getUserIdentifier();  
}
```

```
public enum LoginResponseCode {  
    SUCCESS,  
    FAILURE;  
}
```

Updated Prototype

```
public void testWebServer(WebServer server) {  
    LoginResponse loginResponse = server.login(new  
    LoginRequest());  
  
    // load their profile  
    if (loginResponse.getResponseCode() ==  
    LoginResponseCode.SUCCESS) {  
        // etc  
    }
```

Better for flexibility, but we've sacrificed some usability

Encapsulated Booleans

```
public enum LoginResponseCode {  
    SUCCESS(true),  
    FAILURE(false);  
  
    private boolean success;  
  
    private LoginResponseCode(boolean success) {  
        this.success = success;  
    }  
  
    public boolean success() {  
        return success;  
    }  
}
```

Adding a layer of **indirection** gives us flexibility plus usability

More Updated Prototype

```
public void testWebServer(WebServer server) {  
    LoginResponse loginResponse = server.login(new  
    LoginRequest());  
  
    // load their profile  
    if (loginResponse.getResponseCode().success()) {  
        // etc  
    }
```

That looks better! The 'easy thing' (see if a login happened) is still easy, but now more complicated things (details of the login success/failure) are possible

Encapsulated Booleans

```
public enum LoginResponseCode {  
    SUCCESS(true),  
    FAILURE(false),  
    UNKNOWN_USERNAME(false);  
  
    private boolean success;  
  
    private LoginResponseCode(boolean success) {  
        this.success = success;  
    }  
  
    public boolean success() {  
        return success;  
    }  
}
```

We can add more detailed failure modes in a **backwards-compatible** way

Default Values

"Make the common case easy"

Wherever possible, provide default values in the API (usually **configuration** settings)

- Built-in example use case
- Low cost to start using the API
- Note that changing a default is a **semantic** break; use commonly-useful defaults rather than sentinel values

Ex: the number of threads to use for an execution. Depending on the program, you might use:

- 1 $\Rightarrow$ Multi-threading is an advanced option not generally needed
- 4 $\Rightarrow$ Given how the code is written, more threads will tend to get stuck waiting on locks
- `Runtime.getRuntime().availableProcessors()` $\Rightarrow$ Allows the program to self-tune by default
- -1 $\Rightarrow$ This would be a bad default value

Default Values

How to implement this?

For **cross-network/cross-process** APIs, typically done during deserialization

For **in-process** APIs, typically done by overloading a method:

```
public default int calculateValue(int input) {  
    return calculateValue(input, DEFAULT_NUM_THREADS);  
}  
  
public int calculateValue(int input, int nThreads);
```

Let's Review!

What you should know about API Design that might be on an exam:

Design principles for an API: What the design process is, what concerns to keep in mind, common pitfalls/best practices

Prototyping an API: What this looks like, why it's useful

Types of backwards compatibility: What they are, when to use them

Systems Diagrams: What they are, how to draw them, why to use them

Checkpoint 2: APIs for the Project

Github Issues:

Useful for tracking bugs, feature requests, and (for this assignment) **tasks**

Associate with a **pull request**

Optional, can be good practice with a bug queue

Lightweight, not integrated with github: Asana